Clustering, classification and evaluation

Mostafa H. Chehreghani

Mostafa.chehreghani@gmail.com

Clustering

Albert Bifet (@abifet)

Paris, 18 October 2015
albert.bifet@telecom-paristech.fr

Clustering

Definition

Clustering is the distribution of a set of instances of examples into non-known groups according to some common relations or affinities.

Example

Market segmentation of customers

Example

Social network communities

Clustering

Definition

Given

- a set of instances I
- a number of clusters K
- an objective function $cost(I)$

a clustering algorithm computes an assignment of a cluster for each instance

$$f : I \rightarrow \{1, \ldots, K\}$$

that minimizes the objective function $cost(I)$

Clustering

Definition

Given

- a set of instances I
- a number of clusters K
- an objective function $cost(C, I)$

a clustering algorithm computes a set C of instances with $|C| = K$ that minimizes the objective function

$$cost(C, I) = \sum_{x \in I} d^2(x, C)$$

where

- $d(x, c)$: distance function between x and c
- $d^2(x, C) = min_{c \in C} d^2(x, c)$: distance from x to the nearest point in C

k-means

- 1. Choose k initial centers $C = \{c_1, \ldots, c_k\}$
- 2. while stopping criterion has not been met
  - For $i = 1, \ldots, N$
    - find closest center $c_k \in C$ to each instance p_i
    - assign instance p_i to cluster C_k
  - For $k = 1, \ldots, K$
    - set c_k to be the center of mass of all points in C_i

k-means++

- 1. Choose a initial center c_1
- For $k = 2, \ldots, K$
  - select $c_k = p \in I$ with probability $d^2(p, C)/cost(C, I)$
- 2. while stopping criterion has not been met
  - For $i = 1, \ldots, N$
    - find closest center $c_k \in C$ to each instance p_i
    - assign instance p_i to cluster C_k
  - For $k = 1, \ldots, K$
    - set c_k to be the center of mass of all points in C_i

Performance Measures

Internal Measures

- Cluster Cohesion: Measures how closely related are objects in a cluster
- Cluster Separation: Measure how distinct or well separated a cluster is from other clusters
- Silhouette Coefficient: $1 - a/b$ if $a < b$
  - a = average distance of i to the points in its cluster
  - b = min (average distance of i to points in another cluster)

External Measures

- Rand Measure
- F Measure
- Jaccard
- Purity

Distances

Numeric features

- Euclidean:

$$d(x, y) = ||x - y||_2 = \sum (x_i - y_i)^2$$

- Manhattan distance:

$$d(x, y) = ||x - y||_1 = \sum |x_i - y_i|$$

Density based methods

DBSCAN

- ϵ-neighborhood(p): set of points that are at a distance of p less or equal to ϵ
- Core object: object whose ϵ-neighborhood has an overall weight at least μ
- A point p is *directly density-reachable* from q if
  - p is in ϵ-neighborhood(q)
  - q is a core object
- A point p is *density-reachable* from q if
  - there is a chain of points $p_1, \ldots, p_n$ such that p_{i+1} is directly density-reachable from p_i
- A point p is *density-connected* from q if
  - there is point o such that p and q are density-reachable from o

Density based methods

DBSCAN

- A *cluster* C of points satisfies
  - if $p \in C$ and q is density-reachable from p, then $q \in C$
  - all points $p, q \in C$ are density-connected
- A *cluster* is uniquely determined by any of its core points
- A *cluster* can be obtained
  - choosing an arbitrary core point as a seed
  - retrieve all points that are density-reachable from the seed

Density based methods

DBSCAN

- select an arbitrary point *p*
- retrieve all points density-reachable from *p*
- if *p* is a core point, a cluster is formed
- If p is a border point
  - no points are density-reachable from p
  - DBSCAN visits the next point of the database
- Continue the process until all of the points have been processed

DBSCAN

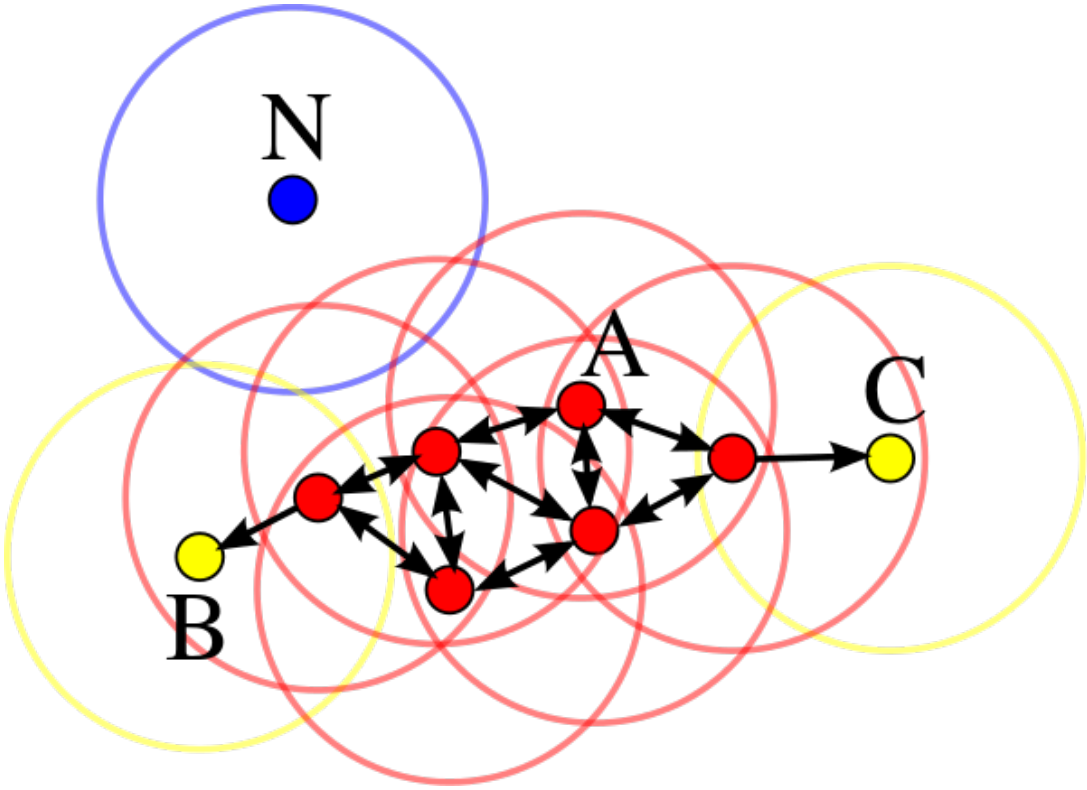

Figure: DBSCAN Point Example with μ=3

BIRCH

Balanced Iterative Reducing and Clustering using Hierarchies

- Clustering Features $CF = (N, LS, SS)$
  - N: number of data points
  - LS: linear sum of the N data points
  - SS: square sum of the N data points
  - Properties:
    - Additivity: $CF_1 + CF_2 = (N_1 + N_2, LS_1 + LS_2, SS_1 + SS_2)$
    - Easy to compute: average inter-cluster distance and average intra-cluster distance
- Uses CF tree
  - Height-balanced tree with two parameters
    - B: branching factor
    - T: radius leaf threshold

BIRCH

Balanced Iterative Reducing and Clustering using Hierarchies

- **Phase 1:** Scan all data and build an initial in-memory CF tree
- **Phase 2:** Condense into desirable range by building a smaller CF tree (optional)
- **Phase 3:** Global clustering
- **Phase 4:** Cluster refining (optional and off line, as requires more passes)